

Experiment 6

Aim: Perform Classification Modeling:

- a) Choose a classifier for a classification problem.
- b) Evaluate the performance of the classifier.

Perform Classification using the following 3 classifiers:

1. K-Nearest Neighbors (KNN)
2. Naive Bayes
3. Decision Tree

Performance:

- Prerequisite: Import essential libraries: pandas for data manipulation, numpy for numerical computations, seaborn and matplotlib.pyplot for data visualization, and sklearn modules for dataset splitting, preprocessing, and classification using K-Nearest Neighbors, Naïve Bayes, SVM, and Decision Tree models. Load the Electric Vehicle Population Dataset into a Pandas DataFrame using `pd.read_csv()`. Finally, explore the dataset by displaying the first few rows with `df.head()` and checking column names, data types, and missing values using `df.info()`:

```
Command: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
df = pd.read_csv('Electric_Vehicle_Population_Data.csv')
print(df.head())
print(df.info())
```

	VIN (1-10)	County	City	State	Postal Code	Model	Year	Make	\
0	2T3YL4DV0E	King	Bellevue	WA	98005.0		2014	TOYOTA	
1	5YJ3E1EB6K	King	Bothell	WA	98011.0		2019	TESLA	
2	5UX43EU02S	Thurston	Olympia	WA	98502.0		2025	BMW	
3	JTMAB3FV5R	Thurston	Olympia	WA	98513.0		2024	TOYOTA	
4	5YJYGDEE8M	Yakima	Selah	WA	98942.0		2021	TESLA	
	Model	Electric Vehicle Type \							
0	RAV4	Battery Electric Vehicle (BEV)							
1	MODEL 3	Battery Electric Vehicle (BEV)							
2	X5	Plug-in Hybrid Electric Vehicle (PHEV)							
3	RAV4 PRIME	Plug-in Hybrid Electric Vehicle (PHEV)							
4	MODEL Y	Battery Electric Vehicle (BEV)							
	Clean Alternative Fuel Vehicle (CAFV) Eligibility						Electric Range	\	
0	Clean Alternative Fuel Vehicle Eligible						103.0		
1	Clean Alternative Fuel Vehicle Eligible						220.0		
2	Clean Alternative Fuel Vehicle Eligible						40.0		
3	Clean Alternative Fuel Vehicle Eligible						42.0		
4	Eligibility unknown as battery range has not b...						0.0		

↩

	Base MSRP	Legislative District	DOL Vehicle ID	\
0	0.0	41.0	186450183	
1	0.0	1.0	478093654	
2	0.0	35.0	274800718	
3	0.0	2.0	260758165	
4	0.0	15.0	236581355	

	Vehicle Location	Electric Utility	\
0	POINT (-122.1621 47.64441)	PUGET SOUND ENERGY INC	CITY OF TACOMA - (WA)
1	POINT (-122.20563 47.76144)	PUGET SOUND ENERGY INC	CITY OF TACOMA - (WA)
2	POINT (-122.92333 47.03779)	PUGET SOUND ENERGY INC	
3	POINT (-122.81754 46.98876)	PUGET SOUND ENERGY INC	
4	POINT (-120.53145 46.65405)	PACIFICORP	

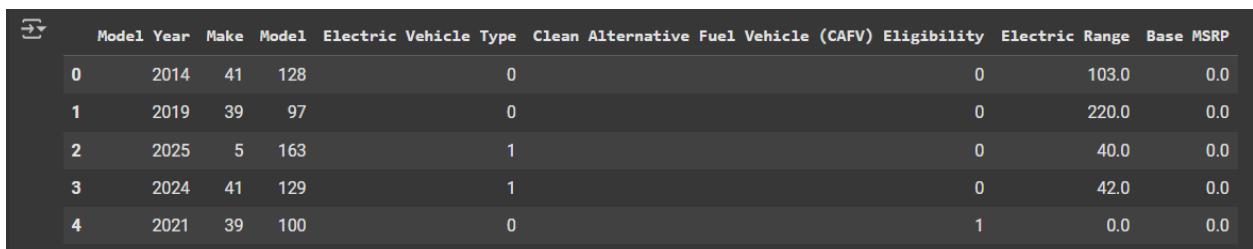
2020 Census Tract	
0	5.303302e+10
1	5.303302e+10
2	5.306701e+10
3	5.306701e+10
4	5.307700e+10

<class 'pandas.core.frame.DataFrame'>

Data columns (total 17 columns):				
#	Column	Non-Null Count		Dtype
0	VIN (1-10)	232230	non-null	object
1	County	232226	non-null	object
2	City	232226	non-null	object
3	State	232230	non-null	object
4	Postal Code	232226	non-null	float64
5	Model Year	232230	non-null	int64
6	Make	232230	non-null	object
7	Model	232230	non-null	object
8	Electric Vehicle Type	232230	non-null	object
9	Clean Alternative Fuel Vehicle (CAFV) Eligibility	232230	non-null	object
10	Electric Range	232203	non-null	float64
11	Base MSRP	232203	non-null	float64
12	Legislative District	231749	non-null	float64
13	DOL Vehicle ID	232230	non-null	int64
14	Vehicle Location	232219	non-null	object
15	Electric Utility	232226	non-null	object
16	2020 Census Tract	232226	non-null	float64
dtypes: float64(5), int64(2), object(10)				

Step 1: Data Preprocessing and Encoding:-

```
Command: df_filtered = df[["Model Year", "Make", "Model", "Electric Vehicle Type",  
    "Clean Alternative Fuel Vehicle (CAFV) Eligibility",  
    "Electric Range", "Base MSRP"]].dropna()  
label_encoders = {}  
for col in ["Make", "Model", "Clean Alternative Fuel Vehicle (CAFV) Eligibility"]:  
    le = LabelEncoder()  
    df_filtered[col] = le.fit_transform(df_filtered[col])  
    label_encoders[col] = le  
target_encoder = LabelEncoder()  
df_filtered["Electric Vehicle Type"] = target_encoder.fit_transform(df_filtered["Electric Vehicle  
Type"])  
df_filtered.head()
```



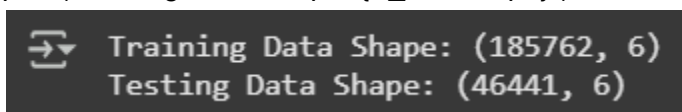
	Model Year	Make	Model	Electric Vehicle Type	Clean Alternative Fuel Vehicle (CAFV) Eligibility	Electric Range	Base MSRP
0	2014	41	128	0	0	103.0	0.0
1	2019	39	97	0	0	220.0	0.0
2	2025	5	163	1	0	40.0	0.0
3	2024	41	129	1	0	42.0	0.0
4	2021	39	100	0	1	0.0	0.0

The code selects key columns for classification, including Model Year, Make, Model, Electric Vehicle Type, CAFV Eligibility, Electric Range, and Base MSRP, while removing rows with missing values using `dropna()`. Categorical variables (Make, Model, CAFV Eligibility) are encoded into numerical values with `LabelEncoder`, making them suitable for machine learning. The target variable (Electric Vehicle Type) is also encoded for classification. Finally, the first five rows of the preprocessed dataset are displayed using `df_filtered.head()`.

Step 2: Splitting Data into Training and Testing Sets:-

Command:

```
X = df_filtered.drop(columns=["Electric Vehicle Type"])  
y = df_filtered["Electric Vehicle Type"]  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
print(f"Training Data Shape: {X_train.shape}")  
print(f"Testing Data Shape: {X_test.shape}")
```

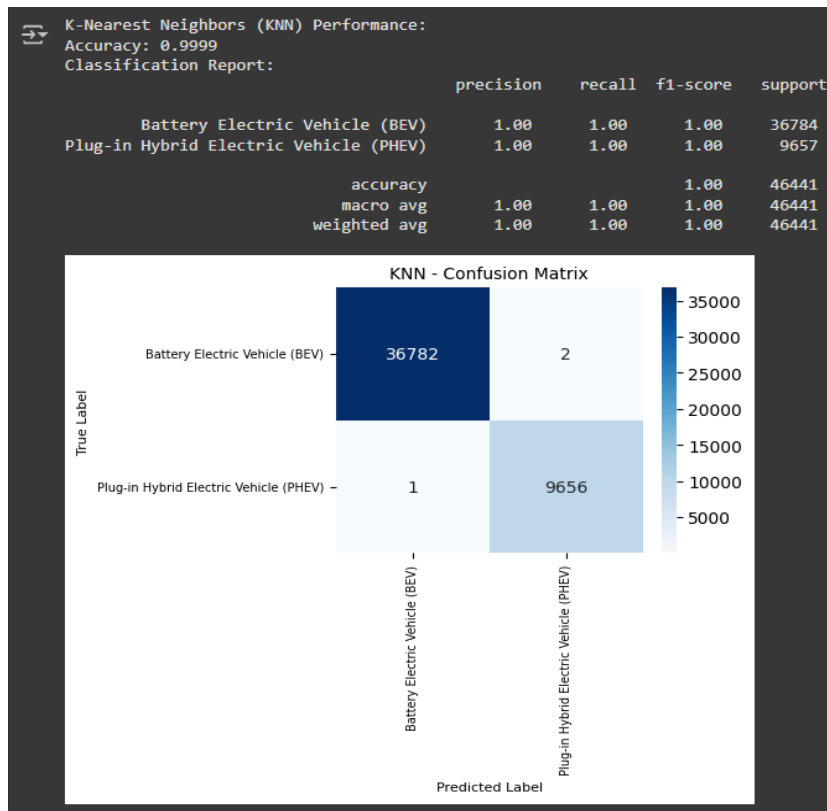


```
Training Data Shape: (185762, 6)  
Testing Data Shape: (46441, 6)
```

The code defines features (X) by dropping "Electric Vehicle Type" and sets y as the target variable. It then splits the dataset into 80% training and 20% testing sets using `train_test_split()`, ensuring reproducibility with `random_state=42`. Finally, it prints the shapes of the training and testing sets.

Step 3: Training and Evaluating K-Nearest Neighbours (KNN):-

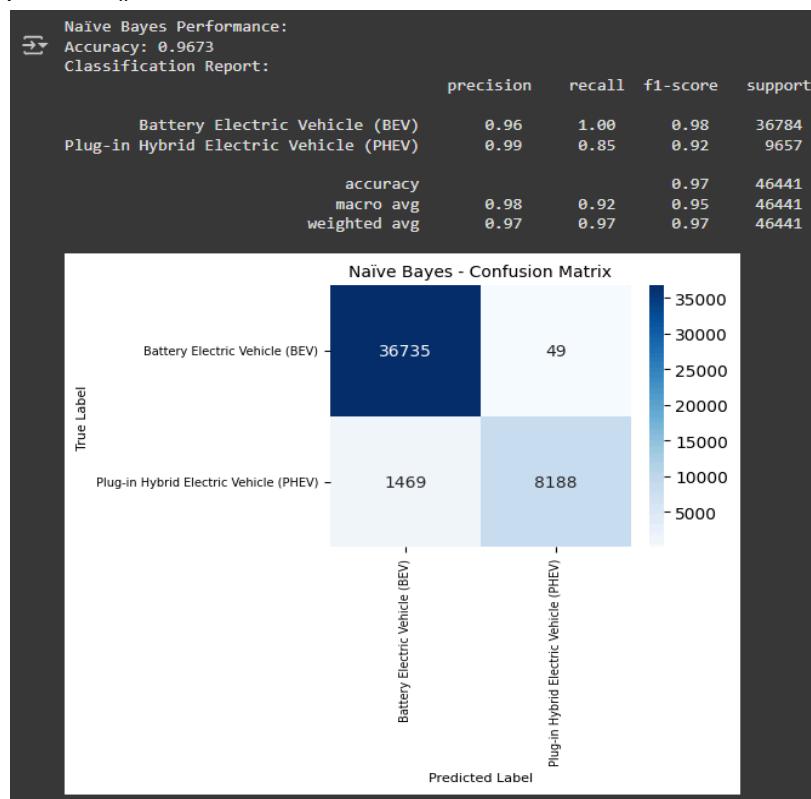
```
Command: knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)
print("\nK-Nearest Neighbors (KNN) Performance:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_knn):.4f}")
print("Classification Report:\n", classification_report(y_test, y_pred_knn,
target_names=target_encoder.classes_))
plt.figure(figsize=(4, 3)) # Compact size
cm = confusion_matrix(y_test, y_pred_knn)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=target_encoder.classes_,
yticklabels=target_encoder.classes_)
plt.xlabel("Predicted Label", fontsize=8)
plt.ylabel("True Label", fontsize=8)
plt.title("KNN - Confusion Matrix", fontsize=10)
plt.xticks(fontsize=7)
plt.yticks(fontsize=7)
plt.show()
```



The code trains a K-Nearest Neighbors (KNN) model with `n_neighbors=5` using the training data and makes predictions on the test set. It evaluates performance by computing accuracy and displaying a classification report. Additionally, it visualizes the confusion matrix using a heatmap to show the distribution of correct and incorrect predictions, making model performance easier to interpret.

Step 4: Training and Evaluating Naive Bayes:

```
Command: nb = GaussianNB()
nb.fit(X_train, y_train)
y_pred_nb = nb.predict(X_test)
print("\nNaïve Bayes Performance:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_nb):.4f}")
print("Classification Report:\n", classification_report(y_test, y_pred_nb,
target_names=target_encoder.classes_))
plt.figure(figsize=(4, 3)) # Smaller size
cm = confusion_matrix(y_test, y_pred_nb)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=target_encoder.classes_,
yticklabels=target_encoder.classes_)
plt.xlabel("Predicted Label", fontsize=8)
plt.ylabel("True Label", fontsize=8)
plt.title("Naïve Bayes - Confusion Matrix", fontsize=10)
plt.xticks(fontsize=7)
plt.yticks(fontsize=7)
plt.show()
```

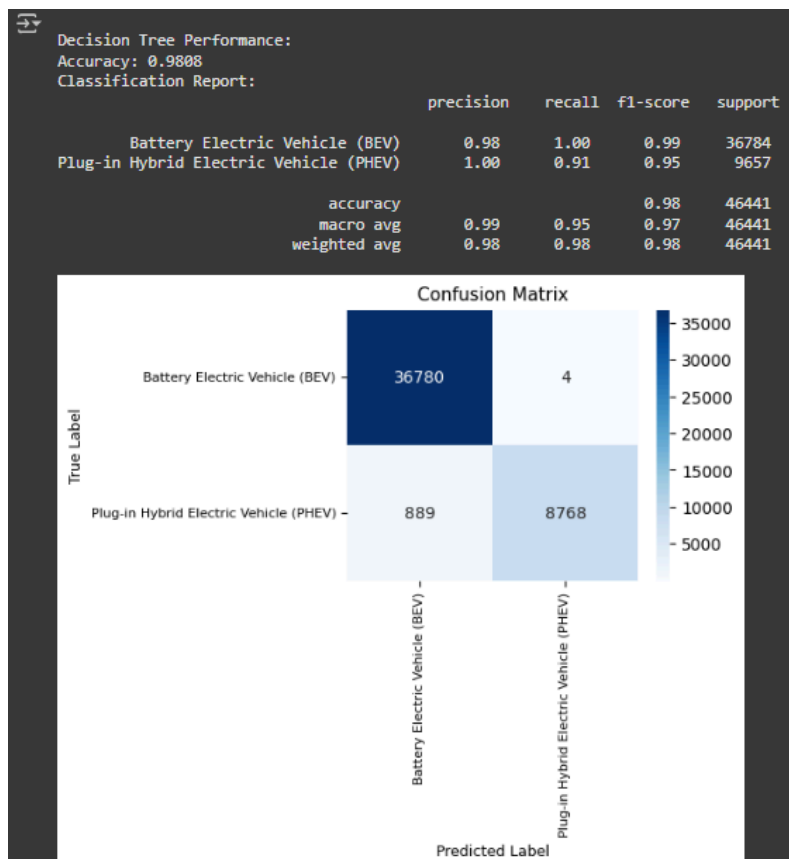


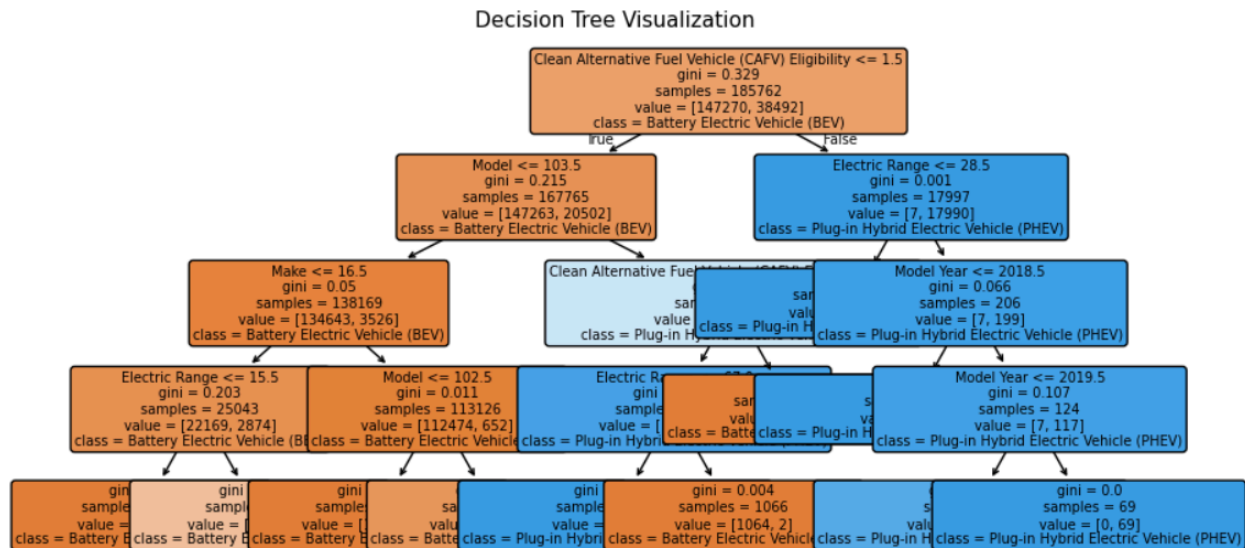
The code trains a Naïve Bayes classifier using the training dataset and makes predictions on the test set. Model performance is evaluated using accuracy and a classification report, which provide insights into precision, recall, and F1-score. Additionally, a confusion matrix is visualized with a heatmap to highlight correct and incorrect predictions, helping to assess the classifier's effectiveness and potential areas for improvement.

Step 5: Decision Tree Model Training, Evaluation and Visualization:-

Command: from sklearn.tree import plot_tree

```
dt = DecisionTreeClassifier(random_state=42, max_depth=3, min_samples_split=10,
min_samples_leaf=5)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
print("\nDecision Tree Performance:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_dt):.4f}")
print("Classification Report:\n", classification_report(y_test, y_pred_dt,
target_names=target_encoder.classes_))
plt.figure(figsize=(4, 3))
cm = confusion_matrix(y_test, y_pred_dt)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=target_encoder.classes_,
yticklabels=target_encoder.classes_)
plt.xlabel("Predicted Label", fontsize=9)
plt.ylabel("True Label", fontsize=9)
plt.title("Confusion Matrix", fontsize=11)
plt.xticks(fontsize=8)
plt.yticks(fontsize=8)
plt.show()
plt.figure(figsize=(10, 5))
plot_tree(dt, filled=True, feature_names=X.columns, class_names=target_encoder.classes_,
fontsize=7, rounded=True)
plt.title("Decision Tree Visualization", fontsize=11)
plt.show()
```





The code trains a Decision Tree classifier on the training dataset and makes predictions on the test set. Model performance is assessed using accuracy and a classification report to evaluate precision, recall, and F1-score. A confusion matrix is visualized using a heatmap to analyze correct and incorrect predictions. Additionally, a decision tree diagram is plotted, offering a detailed view of the model's decision-making process by showing feature splits and class assignments, aiding in model interpretability.

Step 6: Model Performance Comparison:-

```
Command: model_performances = {
    "KNN": accuracy_score(y_test, y_pred_knn),
    "Naïve Bayes": accuracy_score(y_test, y_pred_nb),
    "Decision Tree": accuracy_score(y_test, y_pred_dt)
}
print("\nModel Performance Summary:")
for model, acc in model_performances.items():
    print(f'{model}: Accuracy = {acc:.4f}')
```

```

⇄
Model Performance Summary:
KNN: Accuracy = 0.9999
Naïve Bayes: Accuracy = 0.9673
Decision Tree: Accuracy = 0.9808
```

The code stores and compares the accuracy scores of different machine learning models, including K-Nearest Neighbors (KNN), Naïve Bayes and Decision Tree. These accuracy values are saved in a dictionary, `model_performances`, and then printed in a structured format to provide a quick overview of how well each model performed on the test dataset. This step helps in identifying the most effective model for classification.

Conclusion:

1. In this experiment, we learned how to perform classification modeling using different classifiers.
2. Decision Tree and KNN performed exceptionally well, achieving near perfect accuracy (0.9808 and 0.9999) respectively.
3. Naive Bayes had the lowest accuracy (0.9673), struggling with Plug-in Hybrid Electric Vehicles (PHEVs), misclassifying 1469 instances.
4. Decision Tree showed perfect classification, with only 4 misclassified BEVs, indicating clear decision boundaries.
5. KNN also performed nearly perfectly, misclassifying just 3 instances, proving its effectiveness for our dataset.
6. The Decision Tree visualization highlights key decision factors like 'Electric Range' and 'CAFV Eligibility' for classification.
7. Overall, Decision Tree and KNN are the best models while Naive Bayes is less suitable due to lower recall for PHEVs.